

RAM Disk

Importing the Dependencies

```
[1] import numpy as np
import pandas as pd
from sklearn.model_selection import train_test_split
from sklearn.linear_model import LogisticRegression
from sklearn.metrics import accuracy_score

[2] # loading the dataset to a Pandas DataFrame
credit_card_data = pd.read_csv('/content/creditcard.csv')

[3] # first 5 rows of the dataset
credit_card_data.head()
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	V23	V24	V25	V26	V27	V28	Amount	Class
0	0	-1.359807	-0.072781	2.536347	1.378155	-0.338321	0.462388	0.239599	0.098698	0.363787	...	-0.018307	0.277838	-0.110474	0.066928	0.128539	-0.189115	0.133558	-0.021053	149.62	0.0
1	0	1.191857	0.266151	0.166480	0.448154	0.060018	-0.082361	-0.078803	0.085102	-0.255425	...	-0.225775	-0.638672	0.101288	-0.339846	0.167170	0.125895	-0.008983	0.014724	2.69	0.0
2	1	-1.358354	-1.340163	1.773209	0.379780	-0.503198	1.800499	0.791461	0.247676	-1.514654	...	0.247998	0.771679	0.909412	-0.689281	-0.327642	-0.139097	-0.055353	-0.059752	378.66	0.0
3	1	-0.966272	-0.185226	1.792993	-0.863291	-0.010309	1.247203	0.237609	0.377436	-1.387024	...	-0.108300	0.005274	-0.190321	-1.175575	0.647376	-0.221929	0.062723	0.061458	123.50	0.0
4	2	-1.158233	0.877737	1.548718	0.403034	-0.407193	0.095921	0.592941	-0.270533	0.817739	...	-0.009431	0.798278	-0.137458	0.141267	-0.206010	0.502292	0.219422	0.215153	69.99	0.0

5 rows x 31 columns

[4] credit_card_data.tail()

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V21	V22	V23	V24	V25	V26	V27	V28	Amount	Class
97137	66085	0.998571	-0.046655	0.644378	1.114925	-0.665698	-0.627024	-0.002057	0.018110	-0.052156	...	-0.171808	-0.784227	0.095822	0.455048	0.167887	-0.729060	0.004417	0.037612	89.0	0.0
97138	66085	-1.326193	0.549467	1.220272	1.286509	0.473532	-0.681876	-0.249255	0.444731	-0.768583	...	0.088777	0.029885	-0.123943	-0.092548	-0.159851	-0.360097	0.318036	0.007246	3.6	0.0
97139	66086	1.230983	-0.224520	-0.345196	0.212802	1.586953	3.997378	-1.145351	1.068038	0.584379	...	0.067612	0.229977	-0.119921	1.019614	0.667317	-0.226637	0.071064	0.028365	1.0	0.0
97140	66086	1.241193	0.767604	-0.210715	1.297487	0.152102	-1.162435	0.389686	-0.321743	-0.288129	...	-0.036601	0.032307	-0.136263	0.308814	0.738340	-0.331821	0.040823	0.054137	1.0	0.0
97141	66087	0.310485	-2.576074	1.002015	0.011196	-2.280745	0.465648	-0.860224	0.156411	0.087629	...	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN	NaN

5 rows x 31 columns

[5] # dataset informations

credit_card_data.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 97142 entries, 0 to 97141
Data columns (total 31 columns):
#   Column      Non-Null Count  Dtype
---  -
0   Time        97142 non-null  int64
1   V1          97142 non-null  float64
2   V2          97142 non-null  float64
3   V3          97142 non-null  float64
4   V4          97142 non-null  float64
5   V5          97142 non-null  float64
6   V6          97142 non-null  float64
7   V7          97142 non-null  float64
8   V8          97142 non-null  float64
9   V9          97142 non-null  float64
10  V10         97142 non-null  float64
11  V11         97142 non-null  float64
12  V12         97142 non-null  float64
13  V13         97142 non-null  float64
14  V14         97142 non-null  float64
15  V15         97142 non-null  float64
16  V16         97142 non-null  float64
17  V17         97142 non-null  float64
18  V18         97142 non-null  float64
19  V19         97141 non-null  float64
20  V20         97141 non-null  float64
21  V21         97141 non-null  float64
22  V22         97141 non-null  float64
23  V23         97141 non-null  float64
24  V24         97141 non-null  float64
25  V25         97141 non-null  float64
26  V26         97141 non-null  float64
27  V27         97141 non-null  float64
28  V28         97141 non-null  float64
29  Amount      97141 non-null  float64
30  Class       97141 non-null  float64
dtypes: float64(30), int64(1)
memory usage: 23.0 MB
```

```
[6] # checking the number of missing values in each column
credit_card_data.isnull().sum()
```

```
Time      0
V1        0
V2        0
V3        0
V4        0
V5        0
V6        0
V7        0
V8        0
V9        0
V10       0
V11       0
V12       0
V13       0
V14       0
V15       0
V16       0
V17       0
V18       0
V19       1
V20       1
V21       1
V22       1
V23       1
V24       1
V25       1
V26       1
V27       1
V28       1
Amount    1
Class     1
dtype: int64
```

```
[7] # distribution of legit transactions & fraudulent transactions
credit_card_data['Class'].value_counts()
```

```
Class
0.0    96919
1.0     222
Name: count, dtype: int64
```

This Dataset is highly unblanced

0 -> Normal Transaction

1 -> fraudulent transaction

```
[8] # separating the data for analysis
legit = credit_card_data[credit_card_data.Class == 0]
fraud = credit_card_data[credit_card_data.Class == 1]
```

```
[9] print(legit.shape)
print(fraud.shape)
```

```
(96919, 31)
(222, 31)
```

```
[10] # statistical measures of the data
legit.Amount.describe()
```

```
count    96919.000000
mean      98.310270
std       265.983851
min        0.000000
25%       7.580000
50%      26.610000
75%      89.345000
max     19656.530000
Name: Amount, dtype: float64
```

```
[11] fraud.Amount.describe()
```

```
count     222.000000
mean     114.488243
std      255.373074
min        0.000000
25%        1.000000
50%         7.805000
75%        99.990000
max     1809.600000
Name: Amount, dtype: float64
```

```
[12] # compare the values for both transactions
credit_card_data.groupby('Class').mean()
```

	Time	V1	V2	V3	V4	V5	V6	V7	V8	V9	...	V20	V21	V22	V23	V24	V25	V26	V27	V28	Amount
0.0	41730.096658	-0.249732	-0.043534	0.694778	0.151455	-0.269573	0.098435	-0.093862	0.050744	-0.036243	...	0.043675	-0.031999	-0.108137	-0.036618	0.009741	0.131925	0.026549	-0.000700	0.001378	98.310270
1.0	36541.941441	-6.044462	4.134072	-7.932926	4.915738	-4.386432	-1.796113	-6.300490	2.722455	-2.898611	...	0.345305	0.715188	-0.125165	-0.265450	-0.105791	0.205945	0.103589	0.523395	0.037908	114.488243

2 rows x 30 columns

{x}

Concatenating two DataFrames

[14] new_dataset = pd.concat([legit_sample, fraud], axis=0)

[15] new_dataset.head()

Time V1 V2 V3 V4 V5 V6 V7 V8 V9 ... V21 V22 V23 V24 V25 V26 V27 V28 Amount Class

33968 37457 1.177690 -0.103608 0.392021 0.591672 -0.233191 0.120053 -0.180798 0.066232 0.573304 ... -0.296459 -0.701009 0.024720 -0.427385 0.305461 0.346533 -0.011820 0.010740 29.69 0.0

30875 36114 0.891118 -2.063481 0.503965 -1.256899 -1.902577 0.071625 -1.102155 0.096096 -1.830302 ... 0.177594 0.293194 -0.216890 0.012810 0.125049 -0.118184 0.011053 0.059742 282.00 0.0

49041 43888 -0.120126 -2.382475 1.242680 2.394520 -2.241805 0.499088 -0.278874 0.203564 1.784094 ... 0.312191 -0.042700 -0.583453 0.753939 0.249634 -0.306432 -0.013948 0.146500 623.02 0.0

13529 24000 1.207892 -0.098400 0.472470 0.041344 -0.659200 -0.890118 -0.245740 -0.160615 1.667604 ... -0.100598 -0.080898 0.048033 0.421557 0.112275 1.423444 -0.135263 -0.010233 24.99 0.0

9853 14522 -3.685646 -4.034000 2.340454 0.971496 3.988201 -2.461228 -3.146273 0.669894 1.242409 ... 0.399882 -0.010333 0.591191 -0.343474 0.128209 1.204143 -0.190769 0.128554 15.95 0.0

5 rows x 31 columns

[16] new_dataset.tail()

Time V1 V2 V3 V4 V5 V6 V7 V8 V9 ... V21 V22 V23 V24 V25 V26 V27 V28 Amount Class

95534 65358 1.193916 -0.571085 0.742522 -0.014588 -0.624561 0.832162 -0.833350 0.272897 1.169425 ... -0.049502 0.207265 -0.265272 -0.679294 0.511812 1.246604 -0.028671 -0.006112 31.91 1.0

95597 65385 -2.923827 1.524837 -3.018758 3.289291 -5.755542 2.218276 -0.509995 -3.569444 -1.016592 ... -0.511657 -0.122724 -4.288639 0.563797 -0.949451 -0.204532 1.510206 -0.324706 1354.25 1.0

96341 65728 1.227614 -0.668974 -0.271785 -0.589440 -0.604795 -0.350285 -0.486365 -0.010809 -0.794944 ... -0.026055 -0.295255 -0.180459 -0.436539 0.494649 -0.283738 -0.001128 0.035075 98.01 1.0

96789 65936 -3.593476 0.781442 -1.822448 0.605761 -1.194656 -0.517195 -1.722523 0.128890 0.014963 ... 0.351792 0.391249 -0.252875 -0.498042 0.010172 0.909929 -1.478767 0.722673 101.50 1.0

96994 66037 0.286302 1.399345 -1.682503 3.864377 -1.185373 -0.341732 -2.539380 0.768378 -1.547882 ... 0.352456 -0.243678 -0.194079 -0.172201 0.742237 0.127790 0.569731 0.291206 7.53 1.0

5 rows x 31 columns

[17] new_dataset['Class'].value_counts()

Class

0.0 492

1.0 222

Name: count, dtype: int64

[18] new_dataset.groupby('Class').mean()

Time V1 V2 V3 V4 V5 V6 V7 V8 V9 ... V20 V21 V22 V23 V24 V25 V26 V27 V28 Amount

Class

0.0 41432.353659 -0.379484 -0.175702 0.671829 0.172682 -0.272168 0.090532 -0.10443 0.091357 0.002649 ... 0.074892 -0.032420 -0.116489 0.004871 0.020850 0.153029 0.049331 0.002452 0.011798 114.814715

1.0 36541.941441 -6.044462 4.134072 -7.932926 4.915738 -4.386432 -1.796113 -6.30049 2.722455 -2.896811 ... 0.345305 0.715188 -0.125165 -0.265450 -0.105791 0.205945 0.103589 0.523395 0.037908 114.488243

2 rows x 30 columns

[19] X = new_dataset.drop(columns='Class', axis=1)

Y = new_dataset['Class']

[20] print(X)

Time V1 V2 V3 V4 V5 V6 \

33968 37457 1.177690 -0.103608 0.392021 0.591672 -0.233191 0.120053

30875 36114 0.891118 -2.063481 0.503965 -1.256899 -1.902577 0.071625

49041 43888 -0.120126 -2.382475 1.242680 2.394520 -2.241805 0.499088

13529 24000 1.207892 -0.098400 0.472470 0.041344 -0.659200 -0.890118

9853 14522 -3.685646 -4.034000 2.340454 0.971496 3.988201 -2.461228

...

95534 65358 1.193916 -0.571085 0.742522 -0.014588 -0.624561 0.832162

95597 65385 -2.923827 1.524837 -3.018758 3.289291 -5.755542 2.218276

96341 65728 1.227614 -0.668974 -0.271785 -0.589440 -0.604795 -0.350285

96789 65936 -3.593476 0.781442 -1.822448 0.605761 -1.194656 -0.517195

96994 66037 0.286302 1.399345 -1.682503 3.864377 -1.185373 -0.341732

...

33968 -0.180798 0.066232 0.573304 ... -0.078592 -0.296459 -0.701009

30875 -1.102155 0.096096 -1.830302 ... 0.183339 0.177594 0.293194

49041 -0.278874 0.203564 1.784094 ... 0.997603 0.312191 -0.042700

13529 -0.245740 -0.160615 1.667604 ... -0.170898 -0.100598 -0.080898

9853 -3.146273 0.669894 1.242409 ... 1.175229 0.399882 -0.010333

...

95534 -0.833350 0.272897 1.169425 ... 0.062908 -0.049502 0.207265

95597 -0.509995 -3.569444 -1.016592 ... -0.447039 -0.511657 -0.122724

96341 -0.486365 -0.010809 -0.794944 ... 0.273799 -0.026055 -0.295255

96789 -1.722523 0.128890 0.014963 ... -0.478219 0.351792 0.391249

96994 -2.539380 0.768378 -1.547882 ... 0.270360 0.352456 -0.243678

...

33968 0.024720 -0.427385 0.305461 0.346533 -0.011820 0.010740 29.69

30875 -0.216890 0.012810 0.125049 -0.118184 0.011053 0.059742 282.00

49041 -0.583453 0.753939 0.249634 -0.306432 -0.013948 0.146500 623.02

13529 0.048033 0.421557 0.112275 1.423444 -0.135263 -0.010233 24.99

9853 0.591191 -0.343474 0.128209 1.204143 -0.190769 0.128554 15.95

...

95534 -0.265272 -0.679294 0.511812 1.246604 -0.028671 -0.006112 31.91

95597 -4.288639 0.563797 -0.949451 -0.204532 1.510206 -0.324706 1354.25

96341 -0.180459 -0.436539 0.494649 -0.283738 -0.001128 0.035075 98.01

96789 -0.252875 -0.498042 0.010172 0.909929 -1.478767 0.722673 101.50

96994 -0.194079 -0.172201 0.742237 0.127790 0.569731 0.291206 7.53

...

[714 rows x 30 columns]

```
[21] print(Y)
```

```
33968    0.0
30875    0.0
49041    0.0
13529    0.0
9853     0.0
...
95534    1.0
95597    1.0
96341    1.0
96789    1.0
96994    1.0
Name: Class, Length: 714, dtype: float64
```

Split the data into Training data & Testing Data

```
[22] X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size=0.2, stratify=Y, random_state=2)
```

```
[23] print(X.shape, X_train.shape, X_test.shape)
```

```
(714, 30) (571, 30) (143, 30)
```

Model Training

Logistic Regression

```
[24] model = LogisticRegression()
```

```
[25] # training the Logistic Regression Model with Training Data
model.fit(X_train, Y_train)
```

```
/usr/local/lib/python3.10/dist-packages/sklearn/linear_model/_logistic.py:458: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression

```
n_iter_i = _check_optimize_result(
```

```
  LogisticRegression
  LogisticRegression()
```

Model Evaluation

Accuracy Score

```
[26] # accuracy on training data
X_train_prediction = model.predict(X_train)
training_data_accuracy = accuracy_score(X_train_prediction, Y_train)
```

```
[27] print('Accuracy on Training data : ', training_data_accuracy)
```

```
Accuracy on Training data : 0.957968476357268
```

```
[28] # accuracy on test data
X_test_prediction = model.predict(X_test)
test_data_accuracy = accuracy_score(X_test_prediction, Y_test)
```

```
[29] print('Accuracy score on Test Data : ', test_data_accuracy)
```

```
Accuracy score on Test Data : 0.9230769230769231
```